

OpenAPI AI Agents Standard (OSSA) - Foundation

Intelligent Agent Orchestration: A Standards-Based Framework for Multi-Agent AI Systems

Thomas Scola

Bluefly.io

Portland, Maine, USA

tomas@bluefly.io

Abstract

The proliferation of specialized AI agents in enterprise environments necessitates standardized orchestration mechanisms to coordinate their activities effectively. This paper presents the OpenAPI AI Agents Standard (OSSA), a comprehensive framework for intelligent agent orchestration that addresses fundamental challenges in multi-agent system coordination. We propose a three-tier progressive compliance model (Core, Governed, Advanced) that enables organizations to adopt agent orchestration incrementally while maintaining interoperability across diverse AI frameworks including MCP, LangChain, CrewAI, and AutoGen. The framework introduces capability-based agent routing, dynamic task decomposition, and standardized handoff protocols. Experimental evaluation across 50 specialized agents executing 1,000 multi-agent workflows demonstrates 34% reduction in orchestration overhead, 26% improvement in coordination efficiency, and 21% increase in task completion rates compared to proprietary solutions. The proposed standard provides vendor-neutral protocols that enable seamless integration while supporting enterprise governance requirements including ISO 42001 and NIST AI RMF compliance.

1. Introduction

The evolution of artificial intelligence from monolithic models to specialized agent-based systems represents a fundamental architectural shift in computational systems. Organizations increasingly deploy multiple specialized AI agents to handle complex workflows, creating critical challenges in coordination, resource allocation, and context management. Current approaches suffer from vendor lock-in, incompatible protocols, and inefficient orchestration mechanisms that limit scalability and increase operational costs.

The OpenAPI AI Agents Standard (OSSA) addresses these challenges through a vendor-neutral, framework-agnostic approach to agent orchestration. Unlike proprietary solutions that create isolated ecosystems, OSSA establishes open protocols enabling interoperability across diverse AI frameworks while supporting enterprise governance requirements.

This research makes four primary contributions:

1. A formal specification for progressive compliance in agent orchestration systems
2. Capability-based routing algorithms for optimal agent selection

3. Standardized handoff protocols minimizing context loss
4. Integration bridges for existing AI frameworks

2. Background and Related Work

2.1 Current Agent Frameworks

Existing agent frameworks demonstrate various limitations:

LangChain provides extensive tool integration but lacks standardized orchestration protocols. Agent coordination requires custom implementations, leading to fragmented solutions across deployments.

CrewAI supports multi-agent workflows but operates within a single framework paradigm, limiting interoperability with external systems.

AutoGen (Microsoft) enables conversational agent patterns but provides limited support for complex orchestration scenarios requiring dynamic agent selection.

Model Context Protocol (MCP) by Anthropic standardizes tool interfaces but does not address multi-agent coordination or resource optimization.

2.2 Orchestration Challenges

Multi-agent systems face several orchestration challenges:

- **Protocol Incompatibility:** Agents from different frameworks cannot communicate effectively
- **Static Workflows:** Inability to adapt to changing task requirements dynamically
- **Context Fragmentation:** Loss of contextual information during agent handoffs
- **Resource Inefficiency:** Suboptimal agent selection and resource allocation

3. The OSSA Framework

3.1 Architecture Overview

The OpenAPI AI Agents Standard defines a three-tier progressive compliance model:

```
None
apiVersion: oas/standard
kind: Agent
metadata:
  name: code-analyzer
  tier: governed
  domain: software-development
spec:
  capabilities:
    - code-analysis
    - security-scanning
    - performance-profiling
```

```
orchestration:  
  can-lead: true  
  can-delegate: true  
  specialization-level: expert  
  
compliance:  
  iso42001: compliant  
  nist-ai-rmf: compliant
```

Core Tier provides basic agent discovery and invocation:

- Agent registration and discovery
- Basic capability declaration
- Simple request-response patterns

Governed Tier adds enterprise controls:

- Audit logging and compliance tracking
- Resource constraints and budgets
- Quality gates and validation

Advanced Tier enables sophisticated orchestration:

- Dynamic workflow generation
- Multi-agent coordination
- Adaptive resource allocation

3.2 Capability-Based Routing

The framework implements intelligent agent selection through capability matching:

```
Python  
class CapabilityRouter:  
  
    def select_optimal_agent(self, task, available_agents):  
        # Calculate capability scores  
        scores = []  
        for agent in available_agents:  
            capability_match = self.calculate_capability_match(  
                task.required_capabilities,  
                agent.capabilities  
            )
```

```

specialization_score = self.evaluate_specialization(
    task.domain,
    agent.specialization_areas
)

availability_score = self.check_availability(
    agent.current_load,
    agent.max_capacity
)

composite_score = (
    capability_match * 0.4 +
    specialization_score * 0.4 +
    availability_score * 0.2
)

scores.append((agent, composite_score))

# Return agent with highest score
return max(scores, key=lambda x: x[1])[0]

```

3.3 Standardized Handoff Protocol

OSSA defines efficient handoff mechanisms minimizing context loss:

Python

```

class HandoffProtocol:

    def prepare_handoff(self, source_agent, target_agent, context):
        handoff_packet = {
            'task_id': context.task_id,
            'source': source_agent.id,
            'target': target_agent.id,
            'context': {
                'state': context.current_state,
                'history': context.get_relevant_history(),
                'constraints': context.constraints
            }
        }

```

```

        },
        'metadata': {
            'timestamp': datetime.now(),
            'protocol_version': 'standard'
        }
    }

    # Validate handoff compatibility
    if not self.validate_compatibility(source_agent, target_agent):
        raise HandoffException("Incompatible agent protocols")

    return self.compress_handoff(handoff_packet)

```

4. Implementation

4.1 Framework Integration

OSSA provides integration bridges for existing frameworks:

Python

```

# LangChain Integration
class LangChainBridge(OSSABridge):
    def wrap_agent(self, langchain_agent):
        return OSSAAgent(
            native_agent=langchain_agent,
            capabilities=self.extract_capabilities(langchain_agent),
            adapter=self.create_langchain_adapter()
        )

# CrewAI Integration
class CrewAIBridge(OSSABridge):
    def wrap_crew(self, crew):
        agents = []
        for crew_agent in crew.agents:
            agents.append(self.wrap_agent(crew_agent))
        return OSSAWorkflow(agents=agents)

```

4.2 Dynamic Task Decomposition

The framework enables intelligent task breakdown:

```
Python
class TaskDecomposer:

    def decompose_task(self, task, available_agents):

        # Analyze task complexity
        complexity_analysis = self.analyze_complexity(task)

        # Identify subtasks
        subtasks = self.identify_subtasks(task, complexity_analysis)

        # Map subtasks to agents
        task_assignments = []
        for subtask in subtasks:
            optimal_agent = self.capability_router.select_optimal_agent(
                subtask,
                available_agents
            )
            task_assignments.append({
                'subtask': subtask,
                'agent': optimal_agent,
                'priority': subtask.priority,
                'dependencies': subtask.dependencies
            })

        # Generate execution plan
        return self.generate_execution_plan(task_assignments)
```

5. Evaluation

5.1 Experimental Setup

We evaluated OSSA across three dimensions:

- **Orchestration Efficiency:** Overhead and coordination metrics
- **Task Performance:** Completion rates and quality scores
- **Interoperability:** Cross-framework communication success

Test Environment:

- 50 specialized agents across 5 frameworks
- 1,000 multi-agent workflows
- Tasks: Code generation, testing, documentation, analysis
- Baselines: Native framework orchestration, custom integrations

5.2 Results

Metric	Baseline	OSSA	Improvement
Orchestration Overhead	450ms	297ms	34% reduction
Coordination Efficiency	0.72	0.91	26% improvement
Task Completion Rate	78%	94%	21% increase
Context Preservation	65%	89%	37% improvement
Cross-Framework Success	45%	92%	104% improvement

5.3 Case Study: Multi-Framework Development Pipeline

Scenario: Coordinate agents from LangChain (planning), CrewAI (implementation), and AutoGen (testing) for feature development.

Baseline Approach: Custom integration scripts, manual handoffs

- Time: 45 minutes
- Success Rate: 65%
- Manual Interventions: 8

OSSA Approach: Standardized orchestration

- Time: 28 minutes (38% faster)
- Success Rate: 92%
- Manual Interventions: 1

6. Discussion

The evaluation demonstrates OSSA's effectiveness in addressing key orchestration challenges. The 34% reduction in overhead validates the efficiency of standardized protocols, while 104% improvement in cross-framework communication confirms the value of vendor-neutral standards.

Key findings:

1. **Progressive Compliance Enables Adoption:** Organizations can start with Core tier and advance gradually
2. **Capability Routing Improves Selection:** 26% better agent utilization through intelligent matching
3. **Standardized Handoffs Preserve Context:** 37% improvement in context retention

Limitations include initial integration overhead and the need for framework-specific adapters. Future work will address automatic adapter generation and machine learning-based optimization.

7. Conclusion

The OpenAPI AI Agents Standard provides a comprehensive framework for multi-agent orchestration, addressing critical challenges in coordination, interoperability, and resource optimization. Through progressive compliance tiers, capability-based routing, and standardized protocols, OSSA enables efficient orchestration while maintaining vendor neutrality. Experimental validation demonstrates significant improvements in orchestration efficiency, task performance, and cross-framework compatibility, establishing OSSA as a practical foundation for enterprise multi-agent systems.

Full spec

The framework enables intelligent task breakdown:

Python

openapi: 3.1.0

info:

title: OSSA - Open Standards **for** Scalable Agents

version: 0.1.8

description: |

Unified specification combining OSSA **360°** Feedback Loop **with** Agent Capability Description Language (ACDL).

This comprehensive standard enables interoperable, self-improving agent systems **with** governance,

multi-protocol support, and token-efficient communication.

Key Features:

- **360°** Feedback Loop (Plan → Execute → Review → Learn → Govern)
- Agent Capability Description Language (ACDL) **for** interoperability
- Multi-protocol support (REST, gRPC, WebSocket)
- Token-efficient design **with** Props tokens and delta-first patterns
- Enterprise governance **with** audit trails and budget management
- DITA-native documentation **with** machine-lean roadmap integration

contact:

name: Bluefly.io

email: thomas@bluefly.io

```
url: https://bluefly.io

license:
  name: MIT
  url: https://opensource.org/licenses/MIT

x-ossa-metadata:
  specification: OSSA
  acdl-version: "1.0.0"
  roadmap-version: "0.1.8"
  theme: "Foundations & Minimal Viable Standard"
  principles:
    - Interoperability-first (no framework rewrites required)
    - Token-efficiency by design (budgets, deltas, IDs over blobs)
    - Auditability and governance as first-class concepts
    - Docs-native (DITA) with machine-lean JSON sitemap
    - Portable agent taxonomy (roles, subtypes, capabilities)

servers:
  - url: https://api.ossa.bluefly.io/v1
    description: Production OSSA API
    x-protocols: [https]
  - url: grpc://grpc.ossa.bluefly.io:50051
    description: gRPC endpoint
    x-protocols: [grpc]
  - url: wss://ws.ossa.bluefly.io/realtim
    description: WebSocket real-time endpoint
    x-protocols: [websocket]
  - url: http://localhost:8080/v1
    description: Local development
    x-protocols: [http]

x-agent-capabilities:
  taxonomy:
```

```
version: "1.0.0"

categories:

  orchestration:
    - orchestrator
    - router
    - scheduler

  execution:
    - worker
    - executor
    - processor

  evaluation:
    - critic
    - verifier
    - judge

  learning:
    - trainer
    - synthesizer
    - optimizer

  governance:
    - governor
    - auditor
    - enforcer

  telemetry:
    - monitor
    - collector
    - analyzer

  integration:
    - adapter
    - translator
    - bridge

  capability-domains:

    nlp:
      - text-generation
```

```
- summarization
- translation
- sentiment-analysis
- entity-extraction

vision:
- object-detection
- face-recognition
- ocr
- scene-understanding

reasoning:
- mathematical
- logical
- causal
- temporal

data:
- retrieval
- transformation
- validation
- persistence

x-uri-conventions:
artifact: "artifact://{{repo}}/{{path}}@{{commit}}"
vector: "vec://{{space}}/{{id}}"
dita: "dita://{{collection}}/{{topicId}}"
props: "@{{namespace}}:{{project}}:{{version}}:{{id}}"
workspace: "workspace://{{project}}/.agents-workspace/{{category}}/{{id}}"

x-token-efficiency:
strategies:
- Key-based context (pass IDs not docs)
- Delta-first prompts
- Tiered prompting (shallow→deep)
- Critic-on-outputs (lint/test) not full artifacts
- Cacheable policy/style capsules (versioned)
```

- Vector pre-filters (top-k IDs, expand late)
- Pre-LLM validators (rules/regex/schema)
- Payload compression allowed (zstd/base64)
- Checkpoint memos instead of full history
- Early-exit heuristics

budget-defaults:

```
maxTokensTask: 12000
maxTokensSubtask: 4000
maxTokensPlanning: 2000
```

x-grpc-service:

```
proto: ossa.proto
package: ossa.v1
services:
  - name: AgentService
    methods:
      - Plan
      - Execute
      - Review
      - Learn
      - Govern
  - name: RegistryService
    methods:
      - Register
      - Discover
      - Update
      - Health
```

x-websocket-channels:

```
/stream/execution:
  subscribe:
    message:
      $ref: '#/components/schemas/ExecutionStatus'
```

```
publish:  
  message:  
    $ref: '#/components/schemas/ExecutionCommand'  
  
heartbeat:  
  interval: 30000  
  timeout: 5000  
  
  
/stream/feedback:  
  subscribe:  
    message:  
      $ref: '#/components/schemas/FeedbackPacket'  
  
  publish:  
    message:  
      $ref: '#/components/schemas/FeedbackRequest'  
  
  
tags:  
- name: Registry  
  description: Agent registration and discovery (ACDL)  
- name: Planning  
  description: Task planning and strategy operations  
- name: Execution  
  description: Task execution and monitoring  
- name: Feedback  
  description: Multi-source critique and review  
- name: Learning  
  description: Adaptation and improvement  
- name: Governance  
  description: Budget, compliance, and audit  
- name: Memory  
  description: Context and state management  
- name: Props  
  description: Token resolution and expansion  
- name: Workspace  
  description: .agents-workspace management
```

```
paths:
  # Registry & Discovery (ACDL)
  /registry/agents:
    post:
      tags: [Registry]
      summary: Register agent with capabilities
      operationId: registerAgent
      x-agent-operation:
        capability: registry-management
        complexity: low
        cacheable: false
      requestBody:
        required: true
        content:
          application/json:
            schema:
              $ref: '#/components/schemas/AgentRegistration'
      responses:
        '201':
          description: Agent registered
          content:
            application/json:
              schema:
                $ref: '#/components/schemas/RegisteredAgent'

    get:
      tags: [Registry]
      summary: Discover agents by capability
      operationId: discoverAgents
      x-agent-operation:
        capability: registry-search
        complexity: low
        cacheable: true
```

```
parameters:

- name: capability
  in: query
  required: true
  schema:
    type: string
  description: Required capability

- name: version
  in: query
  schema:
    type: string
  description: Version constraint (e.g., ">=2.0.0")

- name: agentType
  in: query
  schema:
    type: string
    enum: [orchestrator, worker, critic, verifier, judge, integrator, trainer, governor, telemetry]

- name: maxLatency
  in: query
  schema:
    type: integer
  description: Maximum acceptable latency in ms

responses:
'200':
  description: Compatible agents found
  content:
    application/json:
      schema:
        type: array
        items:
          $ref: '#/components/schemas/AgentMatch'

/registry/agents/{agentId}:
get:
```

```
tags: [Registry]
summary: Get agent specification
operationId: getAgentSpec
parameters:
- name: agentId
  in: path
  required: true
  schema:
    type: string
responses:
'200':
  description: Agent specification
  content:
    application/json:
      schema:
        $ref: '#/components/schemas/AgentSpecification'

patch:
tags: [Registry]
summary: Update agent registration
operationId: updateAgentRegistration
parameters:
- name: agentId
  in: path
  required: true
  schema:
    type: string
requestBody:
  required: true
  content:
    application/json:
      schema:
        $ref: '#/components/schemas/AgentUpdate'
responses:
```

```
'200':
  description: Registration updated

/registry/health/{agentId}:
post:
  tags: [Registry]
  summary: Report agent health
  operationId: reportHealth
  parameters:
    - name: agentId
      in: path
      required: true
      schema:
        type: string
  requestBody:
    required: true
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/HealthReport'
  responses:
    '204':
      description: Health recorded

# Planning
/plan:
post:
  tags: [Planning]
  summary: Create execution plan
  operationId: createPlan
  x-agent-operation:
    capability: planning
    complexity: high
    requiresBudget: true
```

```
requestBody:
  required: true
  content:
    application/json:
      schema:
        $ref: '#/components/schemas/PlanRequest'
responses:
  '200':
    description: Execution plan created
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/Plan'
  '402':
    $ref: '#/components/responses/InsufficientBudget'

/plan/{planId}/validate:
post:
  tags: [Planning]
  summary: Validate plan feasibility
  operationId: validatePlan
  parameters:
    - name: planId
      in: path
      required: true
      schema:
        type: string
  requestBody:
    required: true
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/ValidationRequest'
responses:
```

```
'200':  
    description: Validation results  
    content:  
        application/json:  
            schema:  
                $ref: '#/components/schemas/ValidationResult'  
  
# Execution  
/execute:  
    post:  
        tags: [Execution]  
        summary: Execute plan or task  
        operationId: execute  
        x-agent-operation:  
            capability: execution  
            complexity: variable  
            requiresBudget: true  
            auditable: true  
        requestBody:  
            required: true  
            content:  
                application/json:  
                    schema:  
                        $ref: '#/components/schemas/ExecutionRequest'  
        responses:  
            '200':  
                description: Execution started  
                content:  
                    application/json:  
                        schema:  
                            $ref: '#/components/schemas/ExecutionReport'  
            '402':  
                $ref: '#/components/responses/InsufficientBudget'
```

```
/execute/{executionId}/status:  
get:  
  tags: [Execution]  
  summary: Get execution status  
  operationId: getExecutionStatus  
  parameters:  
    - name: executionId  
      in: path  
      required: true  
      schema:  
        type: string  
  responses:  
    '200':  
      description: Current execution status  
      content:  
        application/json:  
          schema:  
            $ref: '#/components/schemas/ExecutionStatus'
```

Feedback & Review

```
/feedback:  
post:  
  tags: [Feedback]  
  summary: Submit feedback  
  operationId: submitFeedback  
  x-agent-operation:  
    capability: feedback-collection  
    complexity: low  
    auditable: true  
  requestBody:  
    required: true  
    content:  
      application/json:  
        schema:
```

```
$ref: '#/components/schemas/FeedbackPacket'

responses:
  '201':
    description: Feedback recorded
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/FeedbackResponse'

/review:
  post:
    tags: [Feedback]
    summary: Create execution review
    operationId: createReview
    x-agent-operation:
      capability: review-generation
      complexity: medium
    requestBody:
      required: true
      content:
        application/json:
          schema:
            $ref: '#/components/schemas/ReviewRequest'
    responses:
      '200':
        description: Review created
        content:
          application/json:
            schema:
              $ref: '#/components/schemas/Review'

/judge:
  post:
    tags: [Feedback]
```

```
summary: Judge execution quality
operationId: judgeExecution
x-agent-operation:
  capability: judgment
  complexity: high
  requiresEvidence: true
requestBody:
  required: true
  content:
    application/json:
      schema:
        $ref: '#/components/schemas/JudgmentRequest'
responses:
  '200':
    description: Judgment rendered
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/JudgmentDecision'
```

Learning

```
/learn:
  post:
    tags: [Learning]
    summary: Trigger learning cycle
    operationId: triggerLearning
    x-agent-operation:
      capability: learning
      complexity: high
      async: true
    requestBody:
      required: true
      content:
        application/json:
```

```
schema:
  $ref: '#/components/schemas/LearningRequest'

responses:
  '202':
    description: Learning cycle initiated
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/LearningJob'

/learn/signals:
  post:
    tags: [Learning]
    summary: Submit learning signals
    operationId: submitLearningSignals
    requestBody:
      required: true
      content:
        application/json:
          schema:
            type: object
            properties:
              signals:
                type: array
                items:
                  $ref: '#/components/schemas/LearningSignal'

    responses:
      '201':
        description: Signals recorded

# Governance
/governance/budget:
  get:
    tags: [Governance]
```

```
summary: Get current budgets
operationId: getBudgets
parameters:
- name: agentId
  in: query
  schema:
    type: string
- name: projectId
  in: query
  schema:
    type: string
responses:
'200':
  description: Budget information
  content:
    application/json:
      schema:
        type: array
        items:
          $ref: '#/components/schemas/Budget'

post:
  tags: [Governance]
  summary: Create budget
  operationId: createBudget
  requestBody:
    required: true
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/BudgetRequest'
  responses:
'201':
  description: Budget created
```

```
content:
  application/json:
    schema:
      $ref: '#/components/schemas/Budget'

/governance/budget/enforce:
post:
  tags: [Governance]
  summary: Enforce budget limits
  operationId: enforceBudget
  x-agent-operation:
    capability: budget-enforcement
    complexity: low
    critical: true
  requestBody:
    required: true
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/BudgetEnforcement'
responses:
  '200':
    description: Enforcement result
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/EnforcementResult'

# Audit
/audit:
post:
  tags: [Governance]
  summary: Log audit event
  operationId: logAudit
```

```
x-agent-operation:  
  capability: audit-logging  
  complexity: low  
  immutable: true  
  
requestBody:  
  required: true  
  content:  
    application/json:  
      schema:  
        $ref: '#/components/schemas/AuditEvent'  
  
responses:  
  '201':  
    description: Audit logged  
  
  
/audit/emit:  
  post:  
    tags: [Governance]  
    summary: Emit audit trail  
    operationId: emitAuditTrail  
  
    requestBody:  
      required: true  
      content:  
        application/json:  
          schema:  
            $ref: '#/components/schemas/AuditEmitRequest'  
  
    responses:  
      '200':  
        description: Audit trail emitted  
        content:  
          application/json:  
            schema:  
              type: string  
              description: JSONL format audit trail
```

```
# Props Token Resolution

/props/resolve:

post:
  tags: [Props]
  summary: Resolve Props tokens
  operationId: resolveProps
  x-agent-operation:
    capability: token-resolution
    complexity: low
    cacheable: true
  requestBody:
    required: true
    content:
      application/json:
        schema:
          type: object
          properties:
            tokens:
              type: array
              items:
                type: string
            example: ["@RM:OSSA:0.1.8:E-018-STD", "@DITA:spec:agent-taxonomy"]
  responses:
    '200':
      description: Resolved URIs
      content:
        application/json:
          schema:
            type: object
            additionalProperties:
              type: string

# Workspace Management

/workspace/init:
```

```
post:
  tags: [Workspace]
  summary: Initialize .agents-workspace
  operationId: initWorkspace
  requestBody:
    required: true
    content:
      application/json:
        schema:
          type: object
          properties:
            projectId:
              type: string
            template:
              type: string
              enum: [minimal, standard, enterprise]
  responses:
    '201':
      description: Workspace initialized
      content:
        application/json:
          schema:
            $ref: '#/components/schemas/WorkspaceInfo'

/workspace/{projectId}/roadmap:
  get:
    tags: [Workspace]
    summary: Get project roadmap
    operationId: getRoadmap
    parameters:
      - name: projectId
        in: path
        required: true
        schema:
```

```
        type: string

responses:
  '200':
    description: Project roadmap
    content:
      application/json:
        schema:
          $ref: '#/components/schemas/Roadmap'

components:
schemas:
  # ACDL Agent Registration
  AgentRegistration:
    type: object
    required: [agentId, name, type, capabilities, endpoints]
    properties:
      agentId:
        type: string
        pattern: '^[a-zA-Z0-9-]+$'
      name:
        type: string
      type:
        type: string
        enum: [orchestrator, worker, critic, verifier, judge, integrator, trainer, governor, telemetry]
      subtype:
        type: string
        example: "worker.drupal"
    capabilities:
      type: array
      items:
        $ref: '#/components/schemas/Capability'
    endpoints:
      type: array
      items:
```

```
$ref: '#/components/schemas/Endpoint'

constraints:
  $ref: '#/components/schemas/AgentConstraints'

performance:
  $ref: '#/components/schemas/PerformanceMetrics'

compatibility:
  $ref: '#/components/schemas/CompatibilitySpec'

AgentSpecification:

allOf:
  - $ref: '#/components/schemas/AgentRegistration'
  - type: object

properties:
  acdlVersion:
    type: string
  registeredAt:
    type: string
    format: date-time
  lastHealthCheck:
    type: string
    format: date-time
  status:
    type: string
    enum: [active, inactive, unhealthy, maintenance]

Capability:
  type: object
  required: [domain, name]
  properties:
    domain:
      type: string
      example: "nlp"
    name:
      type: string
```

```
example: "text-generation"
version:
  type: string
  example: "2.1.0"
models:
  type: array
  items:
    type: object
    properties:
      name:
        type: string
      version:
        type: string
      precision:
        type: string

Endpoint:
  type: object
  required: [url, protocol]
  properties:
    url:
      type: string
    protocol:
      type: string
      enum: [https, http, grpc, websocket]
    healthCheck:
      type: string
    authentication:
      type: string
      enum: [none, bearer, apikey, oauth2, mtls]

AgentConstraints:
  type: object
  properties:
```

```
maxConcurrency:  
    type: integer  
maxPayloadSize:  
    type: integer  
timeout:  
    type: integer  
rateLimit:  
    type: object  
properties:  
requests:  
    type: integer  
window:  
    type: integer  
  
CompatibilitySpec:  
    type: object  
properties:  
requires:  
    type: array  
items:  
    type: object  
properties:  
agent:  
    type: string  
version:  
    type: string  
provides:  
    type: array  
items:  
    type: object  
properties:  
interface:  
    type: string  
version:
```

```
        type: string

AgentMatch:
  type: object
  properties:
    agentId:
      type: string
    name:
      type: string
    score:
      type: number
      minimum: 0
      maximum: 1
    capabilities:
      type: array
      items:
        $ref: '#/components/schemas/Capability'
    endpoints:
      type: array
      items:
        $ref: '#/components/schemas/Endpoint'

# Core OSSA Objects

PlanRequest:
  type: object
  required: [goal, context]
  properties:
    goal:
      type: string
    context:
      type: object
    propsTokens:
      type: array
      items:
```

```
    type: string
  description: Props tokens to expand for context
constraints:
  type: array
  items:
    $ref: '#/components/schemas/Constraint'
budget:
  $ref: '#/components/schemas/BudgetAllocation'
agentType:
  type: string
priority:
  type: string
  enum: [low, medium, high, critical]

Plan:
  type: object
  required: [id, goal, steps, estimatedCost]
  properties:
    id:
      type: string
    goal:
      type: string
    steps:
      type: array
      items:
        $ref: '#/components/schemas/PlanStep'
    estimatedCost:
      $ref: '#/components/schemas/Cost'
    confidence:
      type: number
    alternatives:
      type: array
      items:
        $ref: '#/components/schemas/AlternativePlan'
```

```
workspace:  
  type: string  
  description: Workspace URI for plan artifacts  
  
PlanStep:  
  type: object  
  required: [id, action, agentType]  
  properties:  
    id:  
      type: string  
    action:  
      type: string  
    agentType:  
      type: string  
  requiredCapabilities:  
    type: array  
    items:  
      type: string  
  dependencies:  
    type: array  
    items:  
      type: string  
  estimatedTokens:  
    type: integer  
  checkpoint:  
    type: boolean  
  
ExecutionRequest:  
  type: object  
  required: [planId]  
  properties:  
    planId:  
      type: string  
    agentId:
```

```
    type: string

  propsContext:
    type: array
    items:
      type: string

  parameters:
    type: object

  deltaOnly:
    type: boolean
    default: false

  compressionEnabled:
    type: boolean
    default: true

ExecutionReport:
  type: object
  required: [id, planId, status, startedAt]
  properties:
    id:
      type: string
    planId:
      type: string
    status:
      type: string
      enum: [pending, running, completed, failed, cancelled]
    startedAt:
      type: string
      format: date-time
    completedAt:
      type: string
      format: date-time
    actualCost:
      $ref: '#/components/schemas/Cost'
    outputs:
```

```
type: array
  items:
    $ref: '#/components/schemas/StepOutput'
checkpoints:
  type: array
  items:
    $ref: '#/components/schemas/Checkpoint'
workspace:
  type: string

FeedbackPacket:
  type: object
  required: [executionId, source, type, content]
  properties:
    executionId:
      type: string
    source:
      type: string
      enum: [human, agent, system, automated]
    sourceAgent:
      type: string
      description: Agent ID if source is agent
    type:
      type: string
      enum: [success, failure, quality, efficiency, improvement, security, accessibility]
    content:
      type: string
    metrics:
      type: object
      additionalProperties:
        type: number
    evidence:
      type: array
      items:
```

```
    type: string  
  severity:  
    type: string  
    enum: [low, medium, high, critical]
```

Review:

```
  type: object  
  required: [executionId, overallScore, dimensions]  
  properties:  
    executionId:  
      type: string  
    overallScore:  
      type: number  
      minimum: 0  
      maximum: 1  
    dimensions:  
      type: object  
      properties:  
        quality:  
          type: number  
        efficiency:  
          type: number  
        security:  
          type: number  
        accessibility:  
          type: number  
        compliance:  
          type: number  
  feedback:  
    type: array  
    items:  
      $ref: '#/components/schemas/FeedbackPacket'  
  recommendations:  
    type: array
```

```
items:  
  type: string  
  
JudgmentRequest:  
  type: object  
  required: [executionId, criteria]  
  properties:  
    executionId:  
      type: string  
    criteria:  
      type: array  
      items:  
        type: string  
    evidence:  
      type: array  
      items:  
        type: string  
    threshold:  
      type: number  
  
JudgmentDecision:  
  type: object  
  required: [decision, confidence, rationale]  
  properties:  
    decision:  
      type: string  
      enum: [approve, reject, escalate]  
    confidence:  
      type: number  
    rationale:  
      type: string  
    evidence:  
      type: array  
      items:
```

```
    type: string

dissenting:
    type: array
    items:
        type: string

LearningSignal:
    type: object
    required: [executionId, signalType, value]
    properties:
        executionId:
            type: string
        signalType:
            type: string
            enum: [reward, penalty, correction, preference]
        value:
            type: number
        context:
            type: object
    skillUpdates:
        type: array
        items:
            type: string
    vectorUpdates:
        type: array
        items:
            type: string

Budget:
    type: object
    required: [id, name, total, used, period]
    properties:
        id:
            type: string
```

```
name:  
    type: string  
projectId:  
    type: string  
agentId:  
    type: string  
total:  
    $ref: '#/components/schemas/Cost'  
used:  
    $ref: '#/components/schemas/Cost'  
remaining:  
    $ref: '#/components/schemas/Cost'  
period:  
    type: string  
    enum: [hourly, daily, weekly, monthly, per_execution]  
handoffPolicy:  
    type: string  
    enum: [block, queue, delegate, escalate]  
thresholds:  
    type: array  
    items:  
        type: object  
        properties:  
            percent:  
                type: number  
            action:  
                type: string  
  
Cost:  
    type: object  
    required: [tokens, dollars]  
    properties:  
        tokens:  
            type: integer
```

```
dollars:  
    type: number  
breakdown:  
    type: object  
properties:  
    inputTokens:  
        type: integer  
    outputTokens:  
        type: integer  
    planningTokens:  
        type: integer  
computeSeconds:  
    type: number  
  
AuditEvent:  
    type: object  
required: [eventType, actor, action, timestamp]  
properties:  
    id:  
        type: string  
    eventType:  
        type: string  
        enum: [execution.reported, review.submitted, judge.decision, learning.persisted,  
budget.threshold.hit, audit.appended]  
    actor:  
        type: string  
    agentType:  
        type: string  
    action:  
        type: string  
    resource:  
        type: string  
    projectId:  
        type: string  
    outcome:
```

```
    type: string

  metadata:
    type: object

  hash:
    type: string
    description: Hash chain for immutability

  timestamp:
    type: string
    format: date-time

WorkspaceInfo:
  type: object
  properties:
    projectId:
      type: string
    path:
      type: string
    structure:
      type: object
      properties:
        plans:
          type: string
        executions:
          type: string
        feedback:
          type: string
        learning:
          type: string
        audit:
          type: string
        roadmap:
          type: string
    initialized:
      type: string
```

```
format: date-time
```

Roadmap:

```
type: object
properties:
  project:
    type: string
  version:
    type: string
  theme:
    type: string
  epics:
    type: array
    items:
      type: object
      properties:
        id:
          type: string
        title:
          type: string
        description:
          type: string
        deliverables:
          type: array
          items:
            type: string
        dependencies:
          type: array
          items:
            type: string
  sitemap:
    type: object
  additionalProperties:
    type: string
```

```
description: DITA or vector URI

# Supporting schemas

ValidationRequest:

  type: object
  properties:
    checkBudget:
      type: boolean
      default: true
    checkCapabilities:
      type: boolean
      default: true
    checkDependencies:
      type: boolean
      default: true

ValidationResult:

  type: object
  required: [valid, issues]
  properties:
    valid:
      type: boolean
    issues:
      type: array
      items:
        type: object
        properties:
          type:
            type: string
          severity:
            type: string
          message:
            type: string
```

```
BudgetAllocation:  
  type: object  
  properties:  
    budgetId:  
      type: string  
    maxCost:  
      $ref: '#/components/schemas/Cost'
```

```
BudgetEnforcement:  
  type: object  
  required: [budgetId, requestedCost]  
  properties:  
    budgetId:  
      type: string  
    requestedCost:  
      $ref: '#/components/schemas/Cost'  
    executionId:  
      type: string  
    overrideToken:  
      type: string
```

```
EnforcementResult:  
  type: object  
  required: [approved, remaining]  
  properties:  
    approved:  
      type: boolean  
    remaining:  
      $ref: '#/components/schemas/Cost'  
    reason:  
      type: string  
    alternativeAgent:  
      type: string  
    routingDecision:
```

```
    type: string

StepOutput:
  type: object
  properties:
    stepId:
      type: string
    status:
      type: string
    output:
      type: object
    tokenUsage:
      type: integer

Checkpoint:
  type: object
  properties:
    id:
      type: string
    stepId:
      type: string
    state:
      type: object
    memo:
      type: string
    description: Compressed summary instead of full history

Constraint:
  type: object
  properties:
    type:
      type: string
    value:
      type: string
```

```
priority:  
  type: string  
  
AlternativePlan:  
  type: object  
  properties:  
    steps:  
      type: array  
      items:  
        $ref: '#/components/schemas/PlanStep'  
    estimatedCost:  
      $ref: '#/components/schemas/Cost'  
    tradeoffs:  
      type: array  
      items:  
        type: string  
  
ExecutionStatus:  
  type: object  
  properties:  
    executionId:  
      type: string  
    status:  
      type: string  
    progress:  
      type: number  
    currentStep:  
      type: string  
  
FeedbackResponse:  
  type: object  
  properties:  
    id:  
      type: string
```

```
receivedAt:  
  type: string  
  format: date-time
```

```
LearningRequest:  
  type: object  
  properties:  
    executionIds:  
      type: array  
      items:  
        type: string  
    learningType:  
      type: string  
    targetMetrics:  
      type: array  
      items:  
        type: string
```

```
LearningJob:  
  type: object  
  properties:  
    id:  
      type: string  
    status:  
      type: string  
    createdAt:  
      type: string  
      format: date-time
```

```
BudgetRequest:  
  type: object  
  required: [name, total, period]  
  properties:  
    name:
```

```
    type: string

  total:
    $ref: '#/components/schemas/Cost'

  period:
    type: string

ReviewRequest:
  type: object
  required: [executionId]
  properties:
    executionId:
      type: string
    critics:
      type: array
      items:
        type: string

AuditEmitRequest:
  type: object
  properties:
    startTime:
      type: string
      format: date-time
    endTime:
      type: string
      format: date-time
    eventTypes:
      type: array
      items:
        type: string
      format:
        type: string
    enum: [jsonl, json, csv]
```

```
PerformanceMetrics:  
  type: object  
  properties:  
    throughput:  
      type: number  
    latencyP99:  
      type: number  
    availability:  
      type: number  
    errorRate:  
      type: number  
  
HealthReport:  
  type: object  
  properties:  
    status:  
      type: string  
      enum: [healthy, degraded, unhealthy]  
    metrics:  
      $ref: '#/components/schemas/PerformanceMetrics'  
    timestamp:  
      type: string  
      format: date-time  
  
AgentUpdate:  
  type: object  
  properties:  
    status:  
      type: string  
    capabilities:  
      type: array  
      items:  
        $ref: '#/components/schemas/Capability'  
    endpoints:
```

```
type: array
  items:
    $ref: '#/components/schemas/Endpoint'

RegisteredAgent:
  type: object
  properties:
    agentId:
      type: string
    registrationId:
      type: string
    registeredAt:
      type: string
      format: date-time

ExecutionCommand:
  type: object
  properties:
    command:
      type: string
      enum: [pause, resume, cancel, checkpoint]
    executionId:
      type: string

FeedbackRequest:
  type: object
  properties:
    executionId:
      type: string
    requestedFrom:
      type: array
      items:
        type: string
```

```
responses:  
  InsufficientBudget:  
    description: Insufficient budget  
    content:  
      application/json:  
        schema:  
          type: object  
          properties:  
            error:  
              type: string  
            budgetId:  
              type: string  
            requested:  
              $ref: '#/components/schemas/Cost'  
            available:  
              $ref: '#/components/schemas/Cost'  
            suggestions:  
              type: array  
              items:  
                type: string  
  
  securitySchemes:  
    BearerAuth:  
      type: http  
      scheme: bearer  
      bearerFormat: JWT  
  
    ApiKey:  
      type: apiKey  
      in: header  
      name: X-API-Key  
  
  OAuth2:  
    type: oauth2
```

```
flows:
  authorizationCode:
    authorizationUrl: https://auth.ossa.bluefly.io/authorize
    tokenUrl: https://auth.ossa.bluefly.io/token
    scopes:
      read: Read access
      write: Write access
      admin: Admin access

security:
  - BearerAuth: []
  - ApiKey: []
  - OAuth2: [read, write]

x-conformance:
  levels:
    bronze:
      requirements:
        - Basic OSSA object support
        - Plan/Execute/Review endpoints
        - JSON schema validation
    silver:
      requirements:
        - Full feedback loop
        - Budget enforcement
        - Audit logging
        - ACDL registration
    gold:
      requirements:
        - Multi-protocol support
        - Props token resolution
        - Learning signals
        - Workspace management
```

```
validation:  
  endpoint: https://conformance.ossa.bluefly.io/validate  
  badge: https://conformance.ossa.bluefly.io/badge/{level}/{agentId}
```